

Machine Learning

PHYS 453

Dr. Daugherty



ABILENE
CHRISTIAN
UNIVERSITY

Intro to ML Theory

- Basic vocab and techniques
- Set up so we can use the tools properly
 - it is very easy to skip vital steps and end up with nonsense
 - will see how this setup looks in python with scikit-learn next
- For today we don't care how the models work. We will start caring very soon.

Theory Introduction

SUPERVISED LEARNING

High-Level Overview

Input: represent one sample as numbers (called **features**) in an efficient way that is hopefully useful for solving the problem

Output: answer (the **target**) depends on our question:

- Classification: what discrete category does the sample belong to?
- Regression: predict a real number
- Reinforcement: optimize a certain outcome

Model: how we calculate the output from the input. Usually has lots of adjustable parameters that need to be set somehow




Figure 1.2 Regression and classification problems. a) This *regression* model takes a vector of numbers that characterize a property and predicts its price. b) This *multivariate regression* model takes the structure of a chemical molecule and predicts its melting and boiling points. c) This *binary classification* model takes a restaurant review and classifies it as either positive or negative. d) This *multiclass classification* problem assigns a snippet of audio to one of N genres. e) A second multiclass classification problem in which the model classifies an image according to which of N possible objects it might contain.

Supervised Learning

- **Supervised Learning:** make decisions based on known examples
- **Features:** the measured data on a sample
 - stored in **2D array X**
 - think of a spreadsheet with columns of feature data, each row is a sample
- **Target:** the right answer
 - stored in **1D array y**
 - our Training Data: many samples where we already know X and y
- **Predicting:** goal is to predict y for unknown samples given only X

Data Example

Iris dataset:

- 150 flowers
 - 4 features based on length and width measurements
 - target is species with 3 options

iris setosa



iris versicolor



iris virginica



	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
...
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2

Features

Target

X.shape

(150, 4)

```
X[:5] # the first 5 flowers
```

```
array([[5.1, 3.5, 1.4, 0.2],  
       [4.9, 3. , 1.4, 0.2],  
       [4.7, 3.2, 1.3, 0.2],  
       [4.6, 3.1, 1.5, 0.2],  
       [5. , 3.6, 1.4, 0.2]])
```

y.shape

(150,)

3


Big Picture

Typical steps in ML project:

- **Getting Data** – in the real world cleaning the data can be really difficult
- **Preprocessing** – getting data ready to use, only part we've discussed so far is test/train split
- **Training** – choosing models and parameters, “fitting” them to data
- **Evaluating** – reporting how well it works
- **Predicting** – now that you've got a working thing, go out and use it!

A Relevant Example

Sorting incoming fish according to species using optical sensing



Salmon




Sea Bass

Features

Set up a camera and take some sample images to extract **features**:


- Length
- Lightness
- Width
- Fins
- Position of the mouth
- etc...



No guarantee ahead of time what will work!

Classification

How well does length do?




Decision Boundary – optimal classification threshold

Any fish with length < 11 = salmon, length > 11 = sea bass

Classification


The length is a poor feature alone!

Select the lightness as a possible feature.




Two Features

Where to place the decision boundary? Should it be linear?




Optimal Decision Boundary?



Overfitting – you nailed the training data, but do poorly on different fish.
Model doesn't *generalize* well.



Generalization

Goldilocks' third choice?



Bias vs Variance Tradeoff

<http://nbviewer.jupyter.org/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/05.03-Hyperparameters-and-Model-Validation.ipynb>



UNDERFITTING:

- High bias: too simple to capture shape of data
- Low variance: not sensitive to statistical noise

OVERFITTING:

- Low bias: very flexible, can make complicated shapes
- High variance: **VERY** sensitive to statistical noise

Evaluation


- **Overfitting** – model trained to peculiarities in the training data rather than general samples
- **Underfitting** – model not able to perform well on any data

Our best check for this is to compare the model's performance on:


- 1) data used for training the model
- 2) data the model has never seen before

This trick is called the **test/train split**



Test/Train Split



- The simplest and most common weapon against overfitting is the **test/train split**
- Train: use this part of the data for training
- Test: do **not** train on this part of the data, check how well our model works on data it hasn't seen yet
- **You can NOT fairly evaluate with data you trained on!**




- The training score is everywhere higher than the validation score. This is generally the case: the model will be a better fit to data it has seen than to data it has not seen.
- Very **low** model complexity: the training data is **under-fit**, poor predictor both for the training data and for any previously unseen data.
- Very **high** model complexity: the training data is **over-fit**, model predicts the training data very well but fails for any previously unseen data.
- For some intermediate value, the validation curve has a maximum at best trade-off



Fitting data with different degree polynomials. The validation curve shows the $n=3$ is optimal.


Test/Train Split

Remember that every problem is different. No way to tell ahead of time how a model will perform!




Linear

- high bias, low complexity
- underfitting



Smooth

- medium bias and complexity
- Fitting well!



Crazy

- low bias, high complexity
- overfitting

Test/Train Split

	Bad Testing Score	Good Testing Score
Bad Training Score	UNDERFITTING High Bias / Low Variance	Generally impossible
Good Training Score	OVERFITTING Low Bias / High Variance	It works!

Summary

- Features X
- Target y
- Underfitting/Fitting/Overfitting
- Bias/Variance
- Test/Train Split