

# Decision Trees


PHYS 453

Dr Daugherty

# Introduction

- 20 questions
- Guess Who?
- <http://www.r2d3.us/visual-intro-to-machine-learning-part-1/>


# Decision Trees



**FIGURE 8.1.** Classification in a basic decision tree proceeds from top to bottom. The questions asked at each node concern a particular property of the pattern, and the downward links correspond to the possible values. Successive nodes are visited until a terminal or leaf node is reached, where the category label is read. Note that the same question, **Size?**, appears in different places in the tree and that different questions can have different numbers of branches. Moreover, different leaf nodes, shown in pink, can be labeled by the same category (e.g., **Apple**). From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

# Binary Decision Trees

yes/no questions only!



**FIGURE 8.2.** A tree with arbitrary branching factor at different nodes can always be represented by a functionally equivalent binary tree—that is, one having branching factor  $B = 2$  throughout, as shown here. By convention the “yes” branch is on the left, the “no” branch on the right. This binary tree contains the same information and implements the same classification as that in Fig. 8.1. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

# **CONCEPTS**

Suppose we have the following five positive examples (the first three are the same as in Example 4.1):

- p1: Length = 3  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many
- p2: Length = 4  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many
- p3: Length = 3  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = few
- p4: Length = 5  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many
- p5: Length = 5  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = few

and the following negatives (the first one is the same as in Example 4.2):

- n1: Length = 5  $\wedge$  Gills = yes  $\wedge$  Beak = yes  $\wedge$  Teeth = many
- n2: Length = 4  $\wedge$  Gills = yes  $\wedge$  Beak = yes  $\wedge$  Teeth = many
- n3: Length = 5  $\wedge$  Gills = yes  $\wedge$  Beak = no  $\wedge$  Teeth = many
- n4: Length = 4  $\wedge$  Gills = yes  $\wedge$  Beak = no  $\wedge$  Teeth = many
- n5: Length = 4  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = few

Can you tell “by eye” where to make the first split?

Suppose we have the following five positive examples (the first three are the same as in [Example 4.1](#)):

p1: Length = 3  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many

p2: Length = 4  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many

p3: Length = 3  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = few

p4: Length = 5  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = many

p5: Length = 5  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = few

and the following negatives (the first one is the same as in [Example 4.2](#)):


n1: Length = 5  $\wedge$  Gills = yes  $\wedge$  Beak = yes  $\wedge$  Teeth = many

n2: Length = 4  $\wedge$  Gills = yes  $\wedge$  Beak = yes  $\wedge$  Teeth = many

n3: Length = 5  $\wedge$  Gills = yes  $\wedge$  Beak = no  $\wedge$  Teeth = many

n4: Length = 4  $\wedge$  Gills = yes  $\wedge$  Beak = no  $\wedge$  Teeth = many


n5: Length = 4  $\wedge$  Gills = no  $\wedge$  Beak = yes  $\wedge$  Teeth = few



A **feature tree** is a tree such that each internal node (the nodes that are not leaves) is labelled with a feature, and each edge emanating from an internal node is labelled with a literal.

The set of literals at a node is called a **split**.

**Each leaf of the tree represents a logical expression** which is the combination of all nodes encountered on the path from root to leaf



---

**Algorithm**  $\text{GrowTree}(D, F)$  – grow a feature tree from training data.

---

**Input** : data  $D$ ; set of features  $F$ .

**Output** : feature tree  $T$  with labelled leaves.

- 1 **if**  $\text{Homogeneous}(D)$  **then return**  $\text{Label}(D)$ ;
  - 2  $S \leftarrow \text{BestSplit}(D, F)$  ; // e.g., **BestSplit-Class** (Algorithm 5.2)
  - 3 split  $D$  into subsets  $D_i$  according to the literals in  $S$ ;
  - 4 **for** each  $i$  **do**
  - 5     **if**  $D_i \neq \emptyset$  **then**  $T_i \leftarrow \text{GrowTree}(D_i, F)$  ;
  - 6     **else**  $T_i$  is a leaf labelled with  $\text{Label}(D)$ ;
  - 7 **end**
  - 8 **return** a tree whose root is labelled with  $S$  and whose children are  $T_i$
-

To find the best split we try all the features and choose the one that minimizes **impurity**

---

**Algorithm** BestSplit-Class( $D, F$ ) – find the best split for a decision tree.

---

**Input** : data  $D$ ; set of features  $F$ .

**Output** : feature  $f$  to split on.

```
1  $I_{\min} \leftarrow 1;$ 
2 for each  $f \in F$  do
3     split  $D$  into subsets  $D_1, \dots, D_l$  according to the values  $v_j$  of  $f$ ;
4     if  $\text{Imp}(\{D_1, \dots, D_l\}) < I_{\min}$  then
5          $I_{\min} \leftarrow \text{Imp}(\{D_1, \dots, D_l\});$ 
6          $f_{\text{best}} \leftarrow f;$ 
7     end
8 end
9 return  $f_{\text{best}}$ 
```

---

# Impurity

Given two classes  $\oplus$  and  $\ominus$  with samples sizes  $n_+$  and  $n_-$

$$p = \frac{n_+}{n_+ + n_-}$$

*proportion/fraction of + class*

We want impurity that is:

- 0 whenever  $p$  is 0 or 1
- same if we swap  $p$  for  $1-p$
- maximum when  $p=1/2$

# Impurity

Minority class:  $\min(p, 1 - p)$

**Gini Index:**  $2p(1 - p)$



*expected error if we labelled all samples randomly*

Entropy:  $-p \log p - (1 - p) \log(1 - p)$

*number of bits to describe classes*

The most common choice is *Gini* or  $\sqrt{\text{Gini}}$

Entropy and Gini index are sensitive to fluctuations in the class distribution,  
 $\sqrt{\text{Gini}}$  isn't.



**(left)** Impurity functions plotted against the empirical probability of the positive class. From the bottom: the relative size of the minority class,  $\min(\dot{p}, 1 - \dot{p})$ ; the Gini index,  $2\dot{p}(1 - \dot{p})$ ; entropy,  $-\dot{p} \log_2 \dot{p} - (1 - \dot{p}) \log_2(1 - \dot{p})$  (divided by 2 so that it reaches its maximum in the same point as the others); and the (rescaled) square root of the Gini index,  $\sqrt{\dot{p}(1 - \dot{p})}$  – notice that this last function describes a semi-circle. **(right)**

Indicating the impurity of a single leaf  $D_j$  as  $\text{Imp}(D_j)$ , the impurity of a set of mutually exclusive leaves  $\{D_1, \dots, D_l\}$  is defined as a weighted average

$$\text{Imp}(\{D_1, \dots, D_l\}) = \sum_{j=1}^l \frac{|D_j|}{|D|} \text{Imp}(D_j)$$

# Impurity

<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

## sklearn.tree.DecisionTreeClassifier

```
class sklearn.tree.DecisionTreeClassifier(criterion='gini', splitter='best', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=None, random_state=None,  
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, class_weight=None, presort=False)
```

[\[source\]](#)

A decision tree classifier.

Read more in the [User Guide](#).

**Parameters:** **criterion** : string, optional (default="gini")

The function to measure the quality of a split. Supported criteria are "gini" for the Gini impurity and "entropy" for the information gain.

# Get used to reading these!

## 1.10.7. Mathematical formulation

Given training vectors  $x_i \in R^n$ ,  $i=1, \dots, l$  and a label vector  $y \in R^l$ , a decision tree recursively partitions the feature space such that the samples with the same labels or similar target values are grouped together.

Let the data at node  $m$  be represented by  $Q_m$  with  $n_m$  samples. For each candidate split  $\theta = (j, t_m)$  consisting of a feature  $j$  and threshold  $t_m$ , partition the data into  $Q_m^{left}(\theta)$  and  $Q_m^{right}(\theta)$  subsets

$$\begin{aligned} Q_m^{left}(\theta) &= \{(x, y) | x_j \leq t_m\} \\ Q_m^{right}(\theta) &= Q_m \setminus Q_m^{left}(\theta) \end{aligned}$$

The quality of a candidate split of node  $m$  is then computed using an impurity function or loss function  $H()$ , the choice of which depends on the task being solved (classification or regression)

$$G(Q_m, \theta) = \frac{n_m^{left}}{n_m} H(Q_m^{left}(\theta)) + \frac{n_m^{right}}{n_m} H(Q_m^{right}(\theta))$$

Select the parameters that minimises the impurity

$$\theta^* = \operatorname{argmin}_{\theta} G(Q_m, \theta)$$

Recurse for subsets  $Q_m^{left}(\theta^*)$  and  $Q_m^{right}(\theta^*)$  until the maximum allowable depth is reached,  $n_m < \min_{samples}$  or  $n_m = 1$ .

# Get used to reading these!

## 1.10.7.1. Classification criteria

If a target is a classification outcome taking on values  $0, 1, \dots, K-1$ , for node  $m$ , let

$$p_{mk} = \frac{1}{n_m} \sum_{y \in Q_m} I(y = k)$$

be the proportion of class  $k$  observations in node  $m$ . If  $m$  is a terminal node, `predict_proba` for this region is set to  $p_{mk}$ . Common measures of impurity are the following.

Gini:

$$H(Q_m) = \sum_k p_{mk}(1 - p_{mk})$$

Log Loss or Entropy:

$$H(Q_m) = - \sum_k p_{mk} \log(p_{mk})$$

To find the next split:

**Gini Index:**  $2p(1 - p)$

- choose a feature and sort training data along it
- calculate impurity splitting between every point
- add Gini values weighted by number of samples:

$$impurity = \frac{n_1 * Gini_1 + n_2 * Gini_2 + \dots}{n_1 + n_2 + \dots}$$

- Choose feature that gives biggest impurity decrease

Length = [3, 4, 5]      [2+, 0-][1+, 3-][2+, 2-]

Gills = [yes, no]      [0+, 4-][5+, 1-]

Beak = [yes, no]      [5+, 3-][0+, 2-]

Teeth = [many, few]      [3+, 4-][2+, 1-]







Length  $2/10 \cdot 2 \cdot (2/2 \cdot 0/2) + 4/10 \cdot 2 \cdot (1/4 \cdot 3/4) + 4/10 \cdot 2 \cdot (2/4 \cdot 2/4) = 0.35$

Gills  $4/10 \cdot 0 + 6/10 \cdot 2 \cdot (5/6 \cdot 1/6) = 0.17;$

Beak  $8/10 \cdot 2 \cdot (5/8 \cdot 3/8) + 2/10 \cdot 0 = 0.38;$

Teeth  $7/10 \cdot 2 \cdot (3/7 \cdot 4/7) + 3/10 \cdot 2 \cdot (2/3 \cdot 1/3) = 0.48.$


### Decision surface of a decision tree using paired features




- setosa
- versicolor
- virginica

[http://scikit-learn.org/stable/auto\\_examples/ensemble/plot\\_voting\\_decision\\_regions.html#sphx-glr-auto-examples-ensemble-plot-voting-decision-regions-py](http://scikit-learn.org/stable/auto_examples/ensemble/plot_voting_decision_regions.html#sphx-glr-auto-examples-ensemble-plot-voting-decision-regions-py)


Decision Tree (depth=4)




KNN (k=7)



Kernel SVM



Soft Voting



# Tree Size

- By default sklearn keeps growing the tree until no more splits are possible:  
**overfitting by default**
- There are options to limit tree size:

## **max\_depth : int, default=None**

The maximum depth of the tree. If None, then nodes are expanded until all leaves are pure or until all leaves contain less than min\_samples\_split samples.

## **min\_samples\_split : int or float, default=2**

The minimum number of samples required to split an internal node:


- If int, then consider `min_samples_split` as the minimum number.
- If float, then `min_samples_split` is a fraction and `ceil(min_samples_split * n_samples)` are the minimum number of samples for each split.

*Changed in version 0.18: Added float values for fractions.*

## **min\_samples\_leaf : int or float, default=1**

The minimum number of samples required to be at a leaf node. A split point at any depth will only be considered if it leaves at least `min_samples_leaf` training samples in each of the left and right branches. This may have the effect of smoothing the model, especially in regression.

- If int, then consider `min_samples_leaf` as the minimum number.
- If float, then `min_samples_leaf` is a fraction and `ceil(min_samples_leaf * n_samples)` are the minimum number of samples for each node.



# Prediction Probability


Classify a new sample by following the decision tree down to a leaf. Then classify the sample as the majority class in the leaf. We even get a probability!

However, if we didn't restrict the depth of our tree and there weren't ties in the training data, we may end up with everything being 100% certain

# **APPLICATIONS**


# Kinect

## How it works





1. Classify each pixel's probability of being each of 32 body parts
2. Determine probabilistic cluster of body configurations consistent with those parts
3. Present the most probable to the user

# Kinect



# Kinect


Training data



# Kinect



**Figure 5. Example inferences.** Synthetic (top row); real (middle); failure modes (bottom). Left column: ground truth for a neutral pose as a reference. In each example we see the depth image, the inferred most likely body part labels, and the joint proposals show as front, right, and top views (overlaid on a depth point cloud). Only the most confident proposal for each joint above a fixed, shared threshold is shown.





**Figure 6. Training parameters vs. classification accuracy.** (a) Number of training images. (b) Depth of trees. (c) Maximum probe offset.

# Random Forest





One common approach to minimize problems is to create a random forest:

- train lots of trees on different random samples of data
- each tree splits on random feature (instead of minimum impurity)
- the whole forest votes on classification
- **VERY** powerful technique that forms basis for many high-performing classifiers (e.g. XGBoost)

<https://xgboost.readthedocs.io/en/stable/tutorials/model.html>



<http://nbviewer.jupyter.org/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/05.08-Random-Forests.ipynb>



Trees can give an artificial stair-step decision boundary.

But, optimizing the features (e.g. PCA) can be a big help! Imagine rotating the coordinate system

**FIGURE 8.5.** If the class of node decisions does not match the form of the training data, a very complicated decision tree will result, as shown at the top. Here decisions are parallel to the axes while in fact the data is better split by boundaries along another direction. If, however, “proper” decision forms are used (here, linear combinations of the features), the tree can be quite simple, as shown at the bottom. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

# Summary

## Pros:

- extremely fast classification
- works with non-numeric data
- expressive

## Cons:

- prone to overfitting, **LIMIT YOUR MAX DEPTH**
- “stair-step” decision boundary